

Carnegie Mellon University

17-644: Applied Deep Learning

Assignment#1 Report

A. Sklearn Traditional Linear Regression Model vs. TensorFlow Single Neuron Model


Since it's the first assignment of the course, I started with creating a virtual conda environment called "applieddeeplearning" and then installing required packages as below:

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
from tensorflow import keras
from tensorflow.keras import layers
```

Then, I created random x values and the y linear regression function with my selected a and b values as such:

```
x = np.random.rand(1000, 1)
a = 3
b = 5
noise = np.random.normal(0, 0.1, size=(1000, 1))
y = a * x + b + noise
```

I decided to add normal noise since it's the most common noise type I could imagine to add. The noise line plotted is below:



Then performed sklearn linear regression and TensorFlow single neural network respectively. Since I had only one neuron, it wasn't a network but a single, linear regression function node. The resulting plots with R^2 values, weight, bias and coefficients are plotted below:

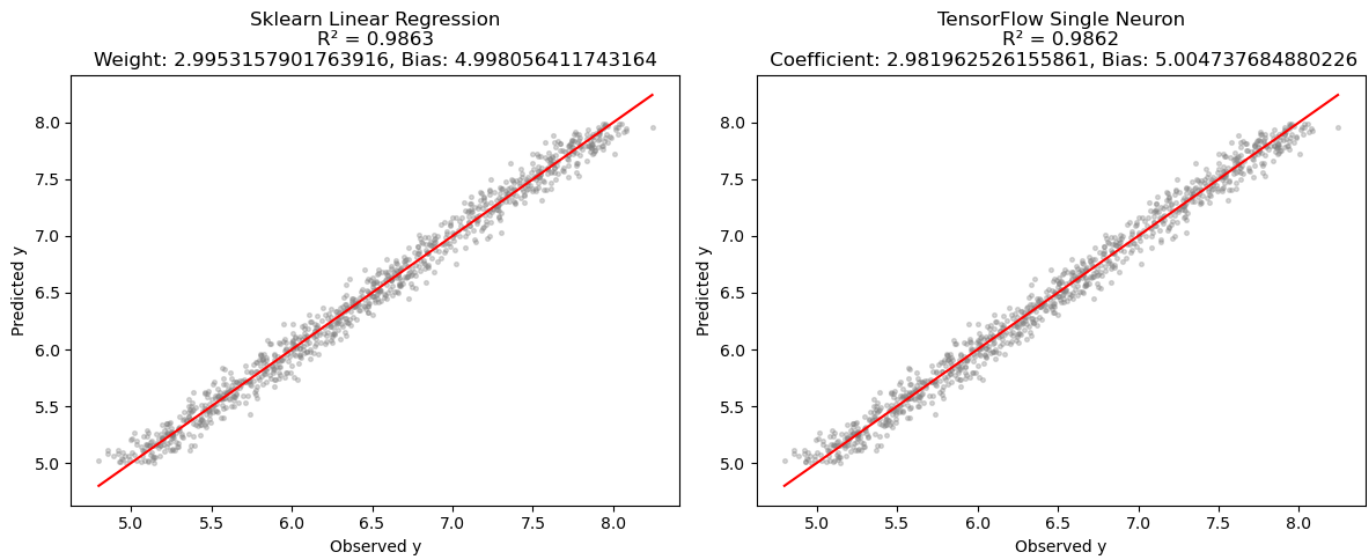


Figure 1 sklearn traditional linear regression model vs. TensorFlow single neuron model.

It wasn't surprising to see TensorFlow performed very close to the traditional sklearn since we didn't use the capacity of deep learning but instead only used one neuron, therefore, launched almost a linear regression algorithm, although in a relatively longer timeframe. As discussed in the lecture, in situations like this, where our dataset isn't that complex, a simple linear regression model could work better than a complicated-look deep learning model.

One important step I had to ensure the TensorFlow code reached the accuracy of the simple linear regression model was to set up the epochs correctly. For this, I had to adjust the epoch value from 200 to 400 as shown in the below code snippet and the plots:

```
tf_model = keras.Sequential([layers.Dense(units=1, input_shape=[1])])
tf_model.compile(optimizer='adam', loss='mean_squared_error')
tf_model.fit(x, y, epochs=400, verbose=0) # played with epochs. From 200 to 400
y_pred = tf_model.predict(x)
r2_tf = r2_score(y, y_pred)
print(f"tensorflow r2_tf: {r2_tf}")
```

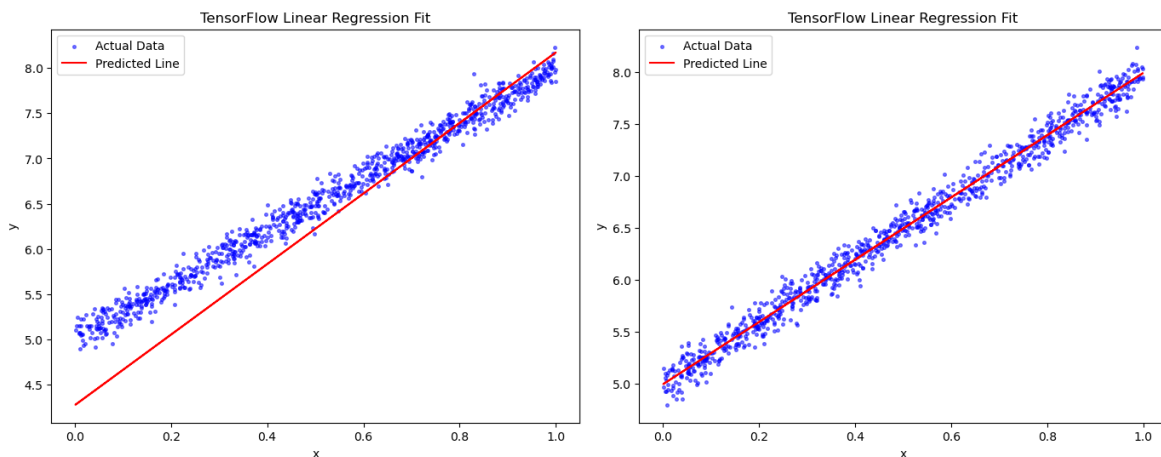


Figure 2 TensorFlow model with epochs = 200 (on the left) and epochs = 400 (on the right)

B. Appendix

The whole code I used to prepare and train the models:

```
x = np.random.rand(1000, 1)
a = 3
b = 5
noise = np.random.normal(0, 0.1, size=(1000, 1))
y = a * x + b + noise
```

```
sk_model = LinearRegression()
sk_model.fit(x, y)
y_pred = sk_model.predict(x)
r2_sk = r2_score(y, y_pred)
print(f"sklearn r2_sk: {r2_sk}")

plt.figure(figsize=(8, 6))
plt.scatter(x, y, color='blue', label='Actual Data', alpha=0.5, s=7)
plt.plot(x, y_pred, color='red', label='Predicted Line')
plt.title('sklearn Linear Regression Fit')
plt.xlabel('x')
plt.ylabel('y')
plt.legend()
plt.show()

print(f"Coefficient: {sk_model.coef_[0][0]}, Bias: {sk_model.intercept_[0]}")
```

```
tf_model = keras.Sequential([layers.Dense(units=1, input_shape=[1])])
tf_model.compile(optimizer='adam', loss='mean_squared_error')
tf_model.fit(x, y, epochs=400, verbose=0) # play with epochs
y_pred = tf_model.predict(x)
r2_tf = r2_score(y, y_pred)
print(f"tensorflow r2_tf: {r2_tf}")

plt.figure(figsize=(8, 6))
plt.scatter(x, y, color='blue', label='Actual Data', alpha=0.5, s=7)
plt.plot(x, y_pred, color='red', label='Predicted Line')
plt.title('TensorFlow Linear Regression Fit')
plt.xlabel('x')
plt.ylabel('y')
plt.legend()
plt.show()

tf_weight, tf_bias = tf_model.layers[0].get_weights()
print(f"Weight: {tf_weight[0][0]}, Bias: {tf_bias[0]}")
```